

# A Brief Introduction to Deep Learning

--Yangyan Li




# A Brief Introduction to Deep Learning

- Artificial Neural Network
- Back-propagation
- Fully Connected Layer
- Convolutional Layer
- Overfitting
- Useful code & examples



# Artificial Neural Network



1. Learning algorithm
2. Neurons
3. Cost function
4. Activation function
5. Weights
6. Models






Live Demo





# Gradient descent



# Neurons are functions

- Let's start with a complex one!

$$f(x, y) = x + y$$

- Given  $x = a, y = b$ , how to update  $x$  *and*  $y$  to make  $f(x, y)$  larger?
- Follow gradient directions!

$$f(x, y) = x + y \quad \rightarrow \quad \frac{\partial f}{\partial x} = 1 \quad \frac{\partial f}{\partial y} = 1$$

$$x = a + 0.01 * 1,$$

$$y = b + 0.01 * 1$$

$$f(x, y): a + b \rightarrow a + b + 0.02$$



# Neurons are functions

- A more complex one!

$$f(x, y) = x * y$$

- Given  $x = a, y = b$ , how to update  $x$  *and*  $y$  to make  $f(x, y)$  larger?
- Follow gradient directions!

$$f(x, y) = xy \quad \rightarrow \quad \frac{\partial f}{\partial x} = y \quad \frac{\partial f}{\partial y} = x$$

$$x = a + 0.01 * b,$$

$$y = b + 0.01 * a$$

$$f(x, y): a * b \rightarrow (a + 0.01 * b)(b + 0.01 * a)$$

$$f(x, y): 4 * (-3) \rightarrow 3.97 * (-2.96)$$




# Back-propagation

- An extremely complex one!



$$f(x, y, z) = (x + y) * z$$

- Let  $q(x, y) = (x + y)$ , then  $f(x, y, z) = q(x, y) * z$

- Chain rule:  $\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$




# Sigmoid $\rightarrow$ ReLU







# Now, serious stuff, a bit...



# Fully Connected Layers

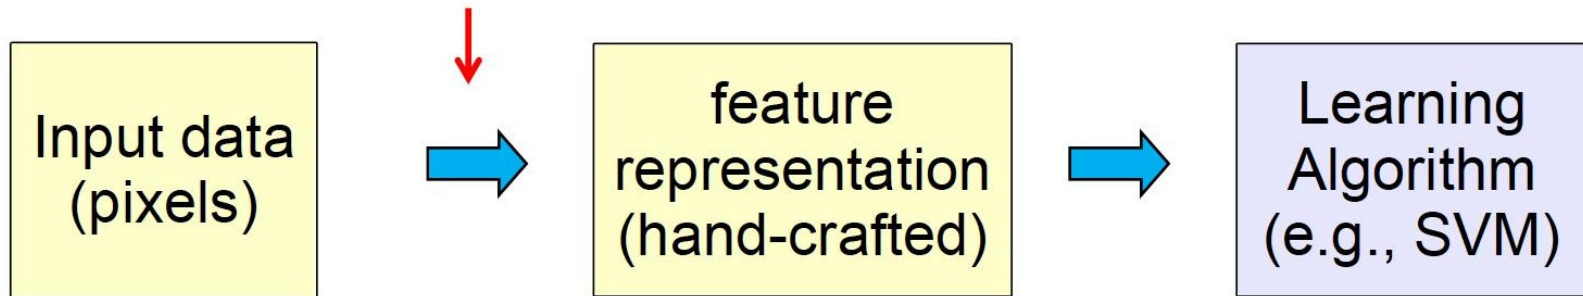


# Convolutional Layers




# Traditional Recognition Approach


Features are not learned



Image



Low-level  
vision features  
(edges, SIFT, HOG, etc.)



Object detection  
/ classification



# Feature Engineering vs. Learning

- Feature engineering is the process of using domain knowledge of the data to create features that make machine learning algorithms work.
- “When working on a machine learning problem, feature engineering is manually designing what the input  $x$ 's should be.”

-- Shayne Miel

- “Coming up with features is difficult, time-consuming, requires expert knowledge.”

--Andrew Ng





With four parameters I can fit an  
elephant, and with five I can make  
him wiggle his trunk.

— *John von Neumann* —


AZ QUOTES







# Under- and Over-fitting examples

Regression:






predictor too inflexible:  
cannot capture pattern




predictor too flexible:  
fits noise in the data



Classification:



# How to detect it in training process?



# Dropout



# A brief history

- McCulloch, Warren S., and Walter Pitts. "A logical calculus of the ideas immanent in nervous activity." *The bulletin of mathematical biophysics* 5.4 (1943): 115-133.
- Rosenblatt, Frank. "The perceptron: a probabilistic model for information storage and organization in the brain." *Psychological review* 65.6 (1958): 386.
- Rumelhart, David E., Geoffrey E. Hinton, and Ronald J. Williams. "Learning representations by back-propagating errors." *Cognitive modeling* 5.3 (1988): 1.
- LeCun, Yann, et al. "Backpropagation applied to handwritten zip code recognition." *Neural computation* 1.4 (1989): 541-551.
- 1993: Nvidia started...
- Hinton, Geoffrey E., Simon Osindero, and Yee-Whye Teh. "A fast learning algorithm for deep belief nets." *Neural computation* 18.7 (2006): 1527-1554.
- Raina, Rajat, Anand Madhavan, and Andrew Y. Ng. "Large-scale deep unsupervised learning using graphics processors." *Proceedings of the 26th annual international conference on machine learning*. ACM, 2009.
- Deng, Jia, et al. "Imagenet: A large-scale hierarchical image database." *Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE Conference on*. IEEE, 2009.
- 2010: "GPUS ARE ONLY UP TO 14 TIMES FASTER THAN CPUS" SAYS INTEL –Nvidia
- Glorot, Xavier, Antoine Bordes, and Yoshua Bengio. "Deep sparse rectifier neural networks." *International Conference on Artificial Intelligence and Statistics*. 2011.
- Hinton, Geoffrey E., et al. "Improving neural networks by preventing co-adaptation of feature detectors." *arXiv preprint arXiv:1207.0580* (2012).
- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." *Advances in neural information processing systems*. 2012.



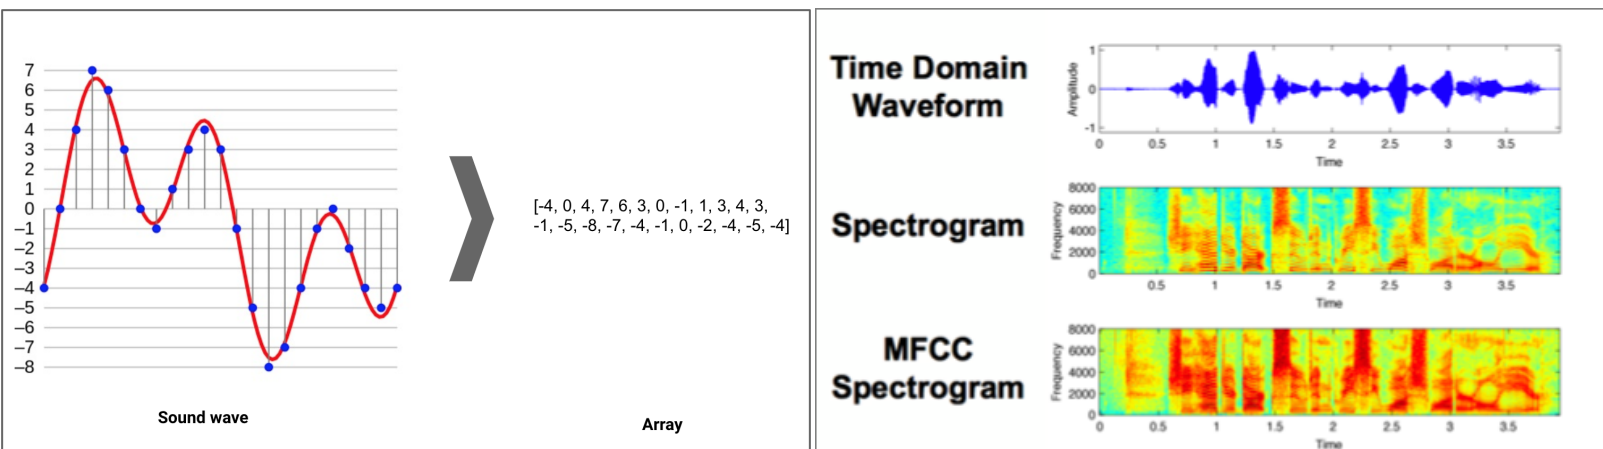
# What are we going to do

- Deep Learning techniques to classify sounds
- Preprocess data, feed it to neural network & train



# Towards a neural network

- For much of the preprocessing we will be able to use Librosa's `load()` function
- Create a visual representation of each of the audio samples which will allow us to identify features for classification, using the same techniques used to classify images with high accuracy
- Images will be the MFCC of the signals
- Useful functions: `librosa.load`, `librosa.feature.mfcc`, `pd.DataFrame`





# Towards a neural network

- Converting the data and labels then splitting the dataset:
  - Using a csv with filename<->digit, load the files and compose an array of [[MFCC,digit] for file in files]
  - Convert into a Panda dataframe
    - `features_df = pd.DataFrame(features, columns=['feature','class_label'])`
  - Encode the classification labels
    - `from sklearn.preprocessing import LabelEncoder`
    - `from keras.utils import to_categorical`
    - `le = LabelEncoder()`
    - `yy = to_categorical(le.fit_transform(y))`
    - `x_train, x_test, y_train, y_test = train_test_split(X, yy, test_size=0.2, random_state = 42)`



# Towards a neural network

- Normalization = make data lie on a scale
- Consider using normalization to speed



# Towards a neural network

- We are going to use CNN

```
model = Sequential()  
model.add(Conv2D(filters=16, kernel_size=2, input_shape=(num_rows, num_columns, num_channels),  
activation='relu'))  
model.add(MaxPooling2D(pool_size=2))  
model.add(Dropout(0.2))  
  
model.add(Conv2D(filters=32, kernel_size=2, activation='relu'))  
model.add(MaxPooling2D(pool_size=2))  
model.add(Dropout(0.2))  
  
model.add(Conv2D(filters=64, kernel_size=2, activation='relu'))  
model.add(MaxPooling2D(pool_size=2))  
model.add(Dropout(0.2))  
  
model.add(Conv2D(filters=128, kernel_size=2, activation='relu'))  
model.add(MaxPooling2D(pool_size=2))  
model.add(Dropout(0.2))  
model.add(GlobalAveragePooling2D())  
  
model.add(Dense(num_labels, activation='softmax'))  
  
model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')
```



“Now this is not the end. It is not even the beginning of the end. But it is, perhaps, the end of the beginning.”

--Winston Churchill

